

You're Losing Organic Traffic to AI.

Here's How to Win It Back.

How answer engines decide which brands to recommend, how to measure whether yours is one of them, and what to check before your competitors do.

Right now, somewhere in Singapore, a buyer is asking ChatGPT: “Who’s the best _____ in Singapore?”

The AI is answering with three or four names. If yours isn’t one of them, this guide is for you.

01 · The Organic Traffic Crisis

Our organic traffic is down 46.8% year on year. We're an SEO agency, and we're telling you this on page one, because the same thing is happening to your site whether you've noticed or not.

4	Organic Search	
	Jan 1 - Jun 10, 2026	
	Jan 1 - Jun 10, 2025	
	% change	-46.83% -46.96%

Between January and June 2026, organic search sessions on our own site fell by nearly half compared with the same period in 2025. Our rankings held. The content didn't get worse. What changed is where the answer lives. For a growing share of searches, it now sits on the results page itself, written by AI, and the click that used to follow never happens.

This isn't a problem unique to one site or one industry. Bain & Company's research, published in February 2025, found that 80% of consumers now rely on AI-written results for at least 40% of their searches, and that the shift has already cut organic web traffic by an estimated 15 to 25%. By Bain's count, roughly 60% of searches now end without the user clicking through to any website at all.

STAT HIGHLIGHTS

60%

of searches now end without a click to any website (Bain & Company, 2025).

15% to 8%

The click rate on results when an AI summary appears (Pew Research, 69K searches).

-34.5%

Clicks for the no.1 ranked page under an AI Overview (Ahrefs, 2025).

The behavioural data tells the same story from the other direction. Pew Research Center analysed almost 69,000 real Google searches from 900 users and found that when an AI summary appears, the click-through rate on traditional results drops from 15% to 8%. The sources cited inside the summary fare even worse: About 1% of users click them. Ahrefs measured the effect on the most valuable position in search and found that the page ranking first loses 34.5% of its clicks when an AI Overview sits above it.

Read those numbers together and the uncomfortable conclusion is hard to avoid. Ranking well no longer guarantees being seen, because the answer engine reads your page, takes what it needs, and serves the conclusion directly. The buyer gets a synthesised answer, usually with a handful of brands recommended by name, and in most cases that's where the journey ends.

The traffic didn't vanish, though. It moved upstream, into the answer itself. Buyers are still researching, comparing and deciding. They've simply stopped doing the legwork on your website and started delegating it to ChatGPT, Gemini, Perplexity and Google's AI surfaces. Which means the question that decides who wins your category has changed. It's no longer only "who ranks?" It's "who gets recommended?"

That question is also why the traffic chart at the start of this section doesn't worry us the way it might appear to. Over the same period our sessions fell, our visibility inside AI answers moved sharply the other way. When someone asks an AI about digital marketing agencies in Singapore, we now appear in 65.7% of answers, at an average position of 1.6, ahead of every competitor we track. The clicks we lost to the answer layer came back as recommendations inside it, and that trade is worth far more than the traffic it replaced: The rest of this guide shows you how to make it too.

02 · Why AI Search Matters More Than The Traffic Numbers Suggest

The instinct is to read all of this as a traffic problem and reach for a traffic solution: Publish more, rank for more, claw the sessions back. That misreads what's happening, because the visits being redistributed by AI are not average visits.

STAT HIGHLIGHT

4.4x

The average AI search visitor converts at 4.4 times the rate of a traditional organic search visitor (Semrush AI search study, 2025).

Semrush's 2025 study of AI search behaviour found that the average visitor arriving from an AI platform converts at 4.4 times the rate of a traditional organic search visitor. Seer Interactive's analysis broke the same pattern down by platform: Referrals from ChatGPT converted at 15.9% against 1.76% for Google organic in the verticals they measured. The volumes are still small as a share of total traffic, but the people in them are disproportionately the ones who buy.

The reason is simple once you picture the journey. By the time someone clicks through from an AI answer, the comparison work is done. The engine has weighed the options, narrowed the shortlist, and presented your brand with a reason attached. The visitor arrives pre-sold and informed, often with their objections already answered. Two decades of marketing effort have gone into manufacturing exactly this kind of visitor. AI answers now produce them by default.

Run that logic in reverse and the cost of invisibility becomes clearer. When an answer engine recommends a competitor instead of you, you haven't lost a click. You've lost a buyer at the final step of their research, and nothing in your analytics will tell you it happened.

There's a second reason this shift deserves attention, and it's better news. Answer engines don't simply copy Google's top 10. Semrush found that close to 90% of ChatGPT's citations come from pages ranking

in position 21 or beyond for related queries. The engines reward content that answers the question cleanly, wherever it sits, which means challenger brands can be recommended ahead of incumbents who have held the rankings for years. Later in this guide we'll show a Singapore telco brand that did exactly that.

For businesses in Singapore, the timing matters. Recommendation patterns in most categories are still being settled, query by query, and the brands doing the structural work now are the ones the engines are learning to trust. That window will not stay open. The rest of this guide covers the two things that determine whether you're inside it: How the engines actually choose, and how to measure where you stand.

03 · How Answer Engines Decide Which Brands To Recommend

Getting recommended by an AI engine looks like luck from the outside. It isn't. It's the output of a retrieval process with observable rules, and once you understand the process, most of the so-called mystery disappears.

How AI Builds An Answer

When someone asks ChatGPT or Google a question, the model doesn't read the whole internet and it doesn't recite a stored answer. It runs a short pipeline. It decides whether the question needs fresh information. If it does, the system retrieves candidate pages, shortlists the ones worth reading, breaks their content into smaller passages (the industry calls these chunks), scores those passages against the question, and assembles a limited set of the best ones into its working context. The answer is then drafted from that selected evidence, with citations pointing back to the passages that survived.

The key idea sits in the middle of that pipeline. The model does not read everything. It builds a small context from selected evidence, and only what makes it into that context can shape the answer. Your brand isn't competing to be crawled. It's competing to be selected, passage by passage.



What We Found When We Studied 342,426 Passages

At The Last Click Summit in 2026, our technical team presented a study built to answer one question: Which passages actually survive retrieval? We analysed 342,426 individual content passages across 12,958 unique URLs cited by AI engines. Five findings matter most for anyone trying to get a brand recommended.

1. **Each engine keeps its own reading list.** Ask ChatGPT and Google AI Overviews the same question and the cited pages barely overlap. Of the 12,958 cited URLs in our study, only 375 were cited by both platforms. AI visibility is not a single ranking you win once. It's a moving probability that differs by platform and by day.
2. **Each platform has its own content appetite.** Google's AI surfaces are comfortable drawing on longer, broader source material. ChatGPT prefers compact, focused pages. A single generic "AI-optimised page" serves neither particularly well; the platform you most need to win on should shape how the content is built.
3. **The opening paragraph is a retrieval gate.** Pages whose first visible passage states directly what the page offers and who it's for were cited across platforms nearly twice as often as pages that warm up slowly: 14.1% against 7.5% in our data. The engines use the opener, together with the title and metadata, to decide whether the page deserves a closer read. Warm-up paragraphs and brand fluff are not neutral. They're a cost.
4. **Medium-sized paragraphs beat dense blocks.** The best-performing passages in our study ran to roughly 2 or 3 focused sentences, with one claim, answer or proof point each. Extra-large blocks of text consistently underperformed. Bigger paragraphs do not produce more AI coverage. Selectable passages do.
5. **Position on the page barely matters. Clarity does.** Across 97,627 passage observations, citation frequency was essentially flat from the top of the page to the bottom. Once a page enters consideration, the engine scans all of it for useful passages. A clear, self-contained answer in your final section is still entirely viable. A vague one at the top is not.

"The opener gets you considered. The chunks get you quoted."


Put together, the findings describe a two-step filter. At page level, your metadata, title, headings and opening paragraph earn you a place in the candidate set. At passage level, the engine then hunts inside the page for self-contained answer blocks it can lift cleanly into its context. Most websites fail at the second step without realising it, because the writing assumes a human who reads top to bottom, with each paragraph leaning on the one before. A passage that only makes sense in sequence makes no sense to a retrieval system.

The Pages AI Reads Are Not Always Yours

There's one more layer, and for some industries it's the decisive one. The page an engine cites when it recommends a brand is often not that brand's page.

Our citation data shows the balance differs sharply by category. In healthcare, medical aesthetics and video production, around 79% of cited pages belong to the brands themselves, so the priority is making your own service pages directly answerable. Education, fitness, wellness and coworking sit in a hybrid

zone, with owned pages accounting for roughly 57 to 70% of citations alongside trusted editorial roundups. In fintech the picture inverts: Owned pages account for closer to 41%, and comparison and listicle content shapes most answers. In categories like F&B catering, competitor pages themselves influence how the engine frames the entire category.



Sourcing habits also differ by platform. Google's AI surfaces draw heavily on community platforms and video, with Semrush's study finding Quora and Reddit the most-cited sites in AI Overviews and YouTube playing a growing role in how products get evaluated. ChatGPT leans on discussion platforms like Reddit alongside business and service websites, which made up half of the links in its answers in the same study.

The practical rule that falls out of all this: Optimise where the engine is already looking. Sometimes that's your service page. Sometimes it's the comparison article that ranks for your category. Sometimes it's a thread where your customers are already talking about you, or conspicuously not. A brand that only tends its own website is optimising a fraction of the surface the engines actually read.

04 · How To Track Your AI Visibility

None of this can be managed on instinct, and very little of it shows up in the tools most businesses already use. Google Analytics will show you a sliver of AI referral traffic, often mislabelled, and it will show you nothing at all about the far larger event: The answers being given about your category, hundreds of times a day, without a single click. Tracking AI visibility means measuring two things, and the distinction between them is where most teams get confused.

Mentions: Are You In The Answer?


A mention is your brand appearing in an AI engine’s answer to a query that matters in your category. It’s the outcome you’re ultimately after, measured across the questions your buyers actually ask. On our own tracking dashboard, when someone asks about a digital marketing agency in Singapore, we’re currently mentioned 65.7% of the time, with a 26.3% share of voice and an average position of 1.6 within the answers. The same view shows every competitor’s numbers beside ours, which is exactly how it should be read: Visibility is relative, and the question is never just “do we appear?” but “who appears instead of us when we don’t?”

AI Search Brand Visibility
Learn and control how AI is talking about your brand

Brand	Visibility		Share of Voice		Avg. Position
 First Page Digital Self	65.7%	-	26.3%	-	1.6
 OOm	51.3%	-	20.5%	-	2.2
 Impossible Marketing	44.1%	-	17.7%	-	2.5
 MediaOne	31.1%	-	12.4%	-	2.9
 Heroes of Digital	24.2%	-	9.7%	-	3.1

Citations: Why You’re In The Answer

A citation is the source the engine actually retrieved to build that answer, the specific page that caused the recommendation. Mentions are the outcome. Citations are the cause. The two can diverge in ways that matter: A brand can be mentioned frequently on the strength of pages it doesn’t own, which is a position of visibility without control, and a brand can hold strong citations on a handful of queries while remaining invisible everywhere else.



Citations in action: Wise’s own pages plus third-party coverage (SingSaver) cause the engine to recommend it by name.

Wise is the clearest public example of citations working as intended. Ask Google’s AI a question like “what platforms can I use to send money internationally with low fees” and Wise comes back as the top recommendation, supported by citations from its own content and from third-party sources like SingSaver. At The Last Click Summit, Wise’s search team described the approach behind it in unglamorous terms: Solid SEO, plus content built to corroborate the long-tail questions users genuinely ask, which earns the citations that drive the recommendation.

They were equally clear about what they didn’t do. Wise doesn’t try to compete on every AI query in its category. It concentrates resources on the questions where it can outshine the competition, and largely concedes the rest.

Three Rules For Measuring It Properly

Topic / Query	Region	Platform	Visibility Score	SOV	Owned Citations
who are the most experienced and trusted seo expe...	SG	All	79.1%	22.5%	105
who are the most experienced and trusted seo cons...	SG	All	78.3%	20.8%	96
who are the best seo companies in singapore	SG	All	96.5%	19.3%	179
who are the best seo agencies in singapore	SG	All	98.0%	20.8%	138

Whatever tool you use, three practices separate measurement from noise. Track ChatGPT and Google’s AI surfaces separately, because as our study showed, they barely share a reading list and a blended score hides more than it reveals.

Measure across consecutive days rather than trusting one snapshot, because AI visibility moves too much for a single check to mean anything.

And report by query cluster, so you can see that you're winning comparison questions but losing recommendation questions, or strong on services and weak on price queries, rather than staring at one average that explains nothing.


05 · What Winning Looks Like

The mechanics in this guide are not theoretical. They're the working method behind NexSEO, our AI search visibility programme, and the results are measurable in exactly the terms this guide has described: Mentions, citations, and share of voice on the queries that matter.

A Singapore Telco: +200% AI visibility

The brand is a challenger competing against larger, better-established incumbents, and when we began, its visibility in AI-generated answers for high-intent queries around roaming and affordability was low. No prior optimisation playbook existed for the category.

The work followed the logic of this guide. We built topical authority around the high-intent themes where the brand could credibly win, created structured, LLM-optimised content designed to be retrieved passage by passage, strengthened the brand and entity signals engines use to establish trust, and tracked visibility across priority prompts continuously, refining the content against what the engines actually selected.



A challenger brand, recommended ahead of incumbents, on the exact queries where buyers decide. And the same approach is compounding elsewhere.

A leading regional bank · +195% · Organic impressions (97K to 287K) while we built AI search readiness across its core banking products through NexSEO's prompt-led content frameworks.

A global logistics leader · +573% · Organic lead volume year on year on its content platform, using answer-optimised content designed for single, high-intent questions, the same passage-level discipline described in section three.

The Answer Engine Checklist

Everything above reduces to a short diagnostic you can run yourself in about 15 minutes, with no tools beyond the AI platforms and a browser. Answer each item honestly: Yes, no, or don't know.

- ☐ Ask ChatGPT, Gemini and Google's AI for the best provider in your category, phrased the way a customer would phrase it. Is your brand in the answer?
- ☐ Ask the same question again over several days and across the platforms. Is your presence consistent, occasional, or absent?
- ☐ Ask one long-tail question a real buyer would ask just before choosing you. Who gets recommended, and which sources does the answer cite?
- ☐ Open your most important service page. Does the first visible paragraph state what you offer, who it's for and the outcome it supports, within the first 2 or 3 sentences?
- ☐ Lift any paragraph from the middle of that page and read it on its own. Does it still make sense, and does it still make a point?
- ☐ Read your H2s and H3s as a list. Do they mirror the questions buyers actually ask, or do they read like slogans?
- ☐ Search the obvious roundup query for your category ("best [your category] in Singapore"). Are you represented on the comparison pages that rank, and represented well?
- ☐ Go back to the citations behind the AI answers from item 1. How many of those sources do you own, influence, or appear on at all?
- ☐ Is anyone in your business tracking brand mentions and citations in AI answers, per platform, over time?
- ☐ Can you name the competitor currently winning the answers you're losing, and the queries they're winning them on?

If you answered "no" or "don't know" to four or more of these, your AI visibility is being settled right now by defaults, competitors and third-party pages, rather than by anything you control. The gap is fixable, but it compounds in whichever direction it's already moving.

Where To Go From Here

The checklist tells you whether you have a problem. It can't tell you precisely where the problem lives, and that's the part worth knowing before you spend anything fixing it.

We run a free AI visibility audit that does exactly that. We put your brand through the same dashboard shown in this guide and map where you're mentioned, where you're cited, which queries your competitors currently own, and which ones are winnable first. You'll leave knowing your actual position in the answers your buyers are reading this week.

[Request your free AI visibility audit →](#)

References & Sources

Bain & Company, "Consumer reliance on AI search results signals new era of marketing", February 2025. [bain.com](https://www.bain.com)

Pew Research Center, analysis of 68,879 Google searches from 900 US adults, published July 2025. [pewresearch.org](https://www.pewresearch.org)

Ahrefs, study of the impact of AI Overviews on click-through rates, 2025. [ahrefs.com](https://www.ahrefs.com)

Semrush, AI search study (visitor value, citation patterns and most-cited domains), 2025. [semrush.com](https://www.semrush.com)

Seer Interactive, LLM referral conversion analysis by platform, 2025. [seerinteractive.com](https://www.seerinteractive.com)

First Page, chunking and citation study: 342,426 content passages across 12,958 cited URLs, presented at The Last Click Summit, 2026.

Client results: campaign data from three First Page clients (a Singapore telco, a regional bank, and a global logistics company; names withheld under client confidentiality), 2025 to 2026.